Management Science

Publication details, including instructions for authors and subscription information:
http://pubsonline.informs.org

Using Charity Performance Metrics as an Excuse Not to Give

Christine L. Exley

To cite this article:
Christine L. Exley (2020) Using Charity Performance Metrics as an Excuse Not to Give. Management Science 66(2):553-563. https://doi.org/10.1287/mnsc.2018.3268

http://pubsonline.informs.org/journal/mnsc

MANAGEMENT SCIENCE
Vol. 66, No. 2, February 2020, pp. 553–563
ISSN 0025-1909 (print), ISSN 1526-5501 (online)

Using Charity Performance Metrics as an Excuse Not to Give

Christine L. Exley[a]

[a] Negotiation, Organizations & Markets Unit, Harvard Business School, Boston, Massachusetts 02163
Contact: clexley@hbs.edu, http://orcid.org/0000-0003-0760-8871 (CLE)

Received: October 19, 2017
Revised: March 28, 2018; September 23, 2018
Accepted: November 21, 2018
Published Online in Articles in Advance: October 24, 2019

https://doi.org/10.1287/mnsc.2018.3268

Abstract. There is an increasing pressure to give more wisely and effectively. There is, relatedly, an increasing focus on charity performance metrics. Through a series of experiments, this paper provides a caution to such a focus. Although information on charity performance metrics may facilitate more effective giving, it may also facilitate the development of excuses not to give. Managers of nonprofit organizations should carefully assess this tension when determining whether and how to provide information on their performance metrics.

History: Accepted by Axel Ockenfels, behavioral economics.
Funding: The author acknowledges funding from the National Science Foundation [Grant SES 1159032] and Harvard Business School.
Supplemental Material: Data files and the online appendix are available at https://doi.org/10.1287/mnsc.2018.3268.

Keywords: charitable giving • prosocial behavior • altruism • excuses • self-serving biases

1. Introduction

Sources ranging from TED Talks to third-party charity evaluators encourage individuals to give wisely.[1] Charity Navigator, a third-party charity evaluator, warns against high overhead costs: "Savvy donors know that the financial health of a charity is a strong indicator of the charity's programmatic performance. . . . [T]he most efficient charities spend 75% or more of their budget on their programs and services."[2] GiveWell more generally encourages caution: "The wrong donation can accomplish nothing. Charities that demonstrably change lives are the exception, not the rule."[3] The Life You Can Save, an organization founded by Peter Singer, echoes this caution by noting that "[n]ot all charities are made the same."[4] Indeed, customized performance reviews for nonprofit organizations are now available via organizations such as ImpactMatters to provide donors with "the analysis needed to make smart, informed funding decisions."[5]

The literature echoes this push from practice toward a focus on charity performance metrics and outcomes.[6] Giving decisions are influenced by the benefit size (Eckel and Grossman 2003, Karlan and List 2007, Meier 2007, Eckel and Grossman 2008, Huck and Rasul 2011, Karlan et al. 2011, Meer 2017), the flexibility in how donations may be used (Gneezy et al. 2014, Batista et al. 2015, Li et al. 2015, Eckel et al. 2017), and effectiveness or efficiency measures such as overhead costs (Gordon et al. 2009, Gneezy et al. 2014, Meer 2014, Metzger and Günther 2015, Yörük 2016, Brown et al. 2017, Coffman 2017, Karlan and Wood 2017).

Although performance metrics clearly influence giving decisions, the value individuals place on performance metrics is often questionable. Among the high-income individuals surveyed in Hope Consulting (2010), only 35% of individuals spend any time researching the performance of charities, even though 85% state that charity performance is very important. Null (2011) and Metzger and Günther (2015) also document a widespread unwillingness to pay to learn information about the impact of potential donations in controlled experiments.[7] Even absent costs associated with learning performance metrics, how individuals respond to the use of donations is often independent of the actual use. Eckel et al. (2017) find that allowing alumni to direct their donations toward their own academic college increases donations even though the directing option is rarely used.[8] Although Gneezy et al. (2014) find reduced giving in response to higher overhead costs, they also find that it is not necessary to reduce overhead costs to encourage more giving. Ensuring potential donors that *others* will cover overhead costs proves effective. The findings in Coffman (2017) further suggest that institutions such as fund-raising campaigns greatly diminish donor sensitivity to overhead costs.

One interpretation from these findings is that there needs to be a greater push in explaining the importance of charity performance metrics. If it is difficult to increase the total amount of charitable giving in the United States—which has hovered at about 2% of the gross domestic product (GDP) for decades—increasing the impact of giving may rely on encouraging

more effective giving.[9] A different interpretation from these findings, which motivates this paper, is that self-serving motives may taint how much individuals appear to value performance metrics. For instance, if participants desire to keep money for themselves and to think of themselves as generous, they may overweight their dislike of less-than-perfect performance metrics as an excuse not to give. They may attribute selfish choices to less-than-perfect performance metrics rather than their own selfishness.[10] Thus, although information on performance metrics may encourage more effective giving, it may also have an unintended consequence of facilitating excuses not to give.

This paper provides evidence consistent with excuse-driven responses to charity performance metrics via two identification strategies. The first strategy relies on comparing how individuals respond to these metrics when self-serving motives are and are not relevant. In the charity–charity treatment, participants make decisions between money going to charities with different performance metrics. Because participants cannot keep money for themselves, and thus self-serving motives are not relevant, the charity–charity treatment captures non-excuse-driven responses to performance metrics. By contrast, in the charity–self treatment, participants make decisions between money going to charities with different performance metrics versus money going to themselves. Because participants can keep money for themselves, and thus self-serving motives are relevant, the charity–self treatment captures excuse-driven responses to performance metrics. The results from a laboratory experiment document a significant difference in response to performance metrics across these two treatments. Although participants respond negatively to lower performance metrics in both treatments, they respond more negatively to performance metrics in the charity–self treatment. Participants appear to overweight their dislike of lower performance metrics as an excuse to keep money for themselves.

The second strategy seeks to vary the ease with which one can develop excuses not to give via a framing manipulation. This approach is similar in spirit to examining whether more selfish behavior arises when a greater degree of flexibility in the decision environment facilitates justifications of more selfish behavior.[11] In the aggregated-information treatment, participants are asked how much they would like to give to charity when their donation will be multiplied by 5, 4.5, 4, 3.5, or 3. In the disaggregated-information treatment, participants are asked how much they would like to give to charity when their donation will be multiplied by 5 and then discounted by a 0%, 10%, 20%, 30%, or 40% "processing fee." Note that participants face the same donation opportunities in both treatments—that is, a donation that is multiplied by 4.5 has an equivalent impact to a donation that is multiplied by 5 and then discounted by 10%. However, because the disaggregated-information treatment explicitly provides "good" information (i.e., donations are multiplied by 5) and "bad" information (donations are then discounted by a processing fee), the disaggregated treatment may facilitate excuse-driven responses because participants may easily overweight the bad information as an excuse not to give. The results from an online experiment are supportive of such excuse-driven responses. Although giving does not decrease as the impact of donations falls in the aggregated-information treatment, giving significantly decreases as the impact of donations falls in the disaggregated-information treatment. When participants can easily overweight their dislike of the processing fee, they appear to do so and keep more money for themselves.

In documenting a novel factor that individuals exploit in a self-serving manner—charity performance metrics—this paper adds to the literature on motivated reasoning.[12] In focusing on how individuals distort their views of unavoidable payoff information from a prosocial action, the most closely related work includes that which documents how individuals exploit ambiguous payoff information (Haisley and Weber 2010) or risky payoff information (Exley 2015). Different than much of the literature on motivated reasoning, the results from the online experiment further show that self-serving responses to unavoidable payoff information can arise even absent uncertainty.[13]

This paper also highlights important policy questions related to the provision of charity performance metrics. Negative responses to performance metrics need not imply the extent to which individuals value better-performing charities. Individuals are prone to responding more negatively to performance metrics when it is self-serving to do so. Caution is warranted when considering what to infer from individuals' responses to charity performance metrics and how to optimize given those responses. To what extent does information on performance metrics encourage more effective giving? How does this compare with the extent that information on performance metrics discourages giving by facilitating the development of excuses not to give? If information on performance metrics is provided, is a focus on aggregate performance metrics sufficient to mitigate excuses not to give? In documenting the importance of these questions and providing some insights, this paper hopes that future work will unpack these questions further.

Section 2 documents excuse-driven responses to performance metrics in a laboratory experiment (Study 1) that employs the first identification strategy. Online Appendix A provides additional and related evidence from two online experiments (Study 1a and

Study 1b), both of which follow a similar design as in Study 1. Section 3 documents excuse-driven responses to performance metrics in an online experiment (Study 2) that employs the second identification strategy. Section 4 concludes.

2. Study 1: A Laboratory Experiment

Study 1 involves data from 50 Stanford University undergraduate students in April 2014 who made a series of binary decisions between money for charities and/or money for themselves. In addition to receiving a $20 completion fee, participants knew that one of their decisions would be randomly selected to count for payment.[14] The design and results for Study 1 are detailed below (and see Online Appendix B.1 for instructions and screenshots).

Three types of charities are involved in Study 1. The first charity type involves three Make-A-Wish Foundation state chapters that vary according to their program expense rates, or percentages of their budgets spent directly on their programs and services (i.e., not spent on overhead costs): the New Hampshire chapter (90%), the Rhode Island chapter (80%), and the Maine chapter (71%).[15] The second charity type involves three Knowledge Is Power Program (KIPP) charter schools that vary according to college matriculation rates among their students who completed the eighth grade: Chicago (92%), Philadelphia (74%), and Denver (61%).[16] The third charity type involves three Bay Area animal shelters that vary according to their live release rates: the San Francisco SPCA (97%), the Humane Society of Silicon Valley (82%), and the San Jose Animal Care and Services (66%).[17]

So that later decisions in the study account for how participants value money for themselves relative to money for charity, the study begins with a normalization procedure that was previously developed in Exley (2015). The normalization procedure determines how participants value money for themselves relative to money for the top-rated charity for each of the three charity types via three normalization price lists. The order in which each participant answers the three normalization price lists is randomly determined, and immediately following each normalization price list, participants complete a "buffer" price list.[18] On each row of a normalization price list, participants choose between (i) $10 for themselves and (ii) some amount for the top-rated charity of type t. Because the amount for the top-rated charity increases by $2 from $0 to $40 as one proceeds down the 21 rows of the price list, the amounts at which participants switch to choosing money for the top-rated charity are informative about how individuals value money for themselves relative to the top-rated charity. More specifically, the switch points imply X_t values such that participants are indifferent between themselves receiving $10 and the top-rated charity of type t receiving X_t.[19] Their X_t values determine the stakes, unbeknownst to participants, involved in the subsequent "valuation" price lists.[20]

The valuation price lists examine how participants value money for themselves versus each lower-rated (second-rated or third-rated) charity. For each lower-rated charity of type t, there are two valuation price lists: one that occurs in the charity–charity treatment and one that occurs in the charity–self treatment. The order of the treatments, as well as the order of charity types within a treatment, is randomly determined. The purpose of the treatments is to examine how individuals' valuations of lower-rated charities differ when self-serving motives are and are not relevant. Although participants cannot choose money for themselves in the charity–charity treatment, and thus non-excuses-driven responses to performance metrics follow, participants can choose money for themselves in the charity–self treatment, and thus (potentially) excuse-driven responses to performance metrics follow.

In the charity–charity treatment, participants choose between (i) $\$X_t$ for a lower-rated charity of type t and (ii) some amount for the top-rated charity of type t that increases in $\$X_t/20$ increments from $0 to $\$X_t$ as one proceeds down the 21 rows of a valuation price list. These decisions imply a charity–charity valuation such that participants are indifferent between the lower-rated charity of type t receiving $\$X_t$ and the top-rated charity of type t receiving some percentage of $\$X_t$. This percentage is called the charity–charity valuation, and the extent to which it falls below 100% indicates the extent to which participants dislike a low-rated charity relative to the corresponding top-rated charity. Equivalently, the charity–charity valuation indicates how often, out of the 21 rows of the valuation price list, participants choose money for a lower-rated charity over money for the corresponding top-rated charity.

In the charity–self treatment, participants choose between (i) $\$X_t$ for a lower-rated charity of type t and (ii) some amount for themselves that increases in $0.50 increments from $0 to $10 as one proceeds down the 21 rows of a valuation price list. Thus, option (i) is the same across both treatments, whereas option (ii) involves the top-rated charity receiving some increasing percentage of $\$X_t$ in the charity–charity treatment but, instead, participants receiving some increasing percentage of $10 for themselves in the charity–self treatment. Given participants are indifferent between $\$X_t$ for the top-rated charity of type t and $10 for themselves, the rows at which participants switch from (i) to (ii) should be the same in both treatments if they are not excuse driven.[21]

If so, participants' charity–charity valuations would be the same as participants' charity–self valuations.[22] If participants are instead less willing to choose $\$X_t$ for a lower-rated charity when self-serving motives are relevant, they may switch more quickly to (ii) in the charity–self treatment than in the charity–charity treatment. More specifically, the extent to which the charity–self valuation is lower than the charity–charity valuation is reflective of the extent of excuse-driven responses to charity performance metrics.

Before turning to the main results of interest from the valuation price lists, a few results from the normalization price lists are informative. Implied X_t values—such that participants are indifferent between themselves receiving \$10 and the top-rated charity of type t receiving $\$X_t$—are unclear for three participants who switched from choosing money for themselves to money for the top-rated charity of type t more than once on a normalization price list. These participants are excluded from the analyses. In 31% of the remaining 141 normalization price lists (three for each of the remaining 47 participants), participants always choose \$10 for themselves, and thus the best estimate of X_t is censored and set to equal the maximum value in the price list: \$40.[23] In all other cases, X_t values are easily inferred by the unique amounts at which participants switch from choosing \$10 for themselves to $\$X_t$ for the top-rated charity t.[24] Figure 1 shows these noncensored X_t values when they are translated into self-to-charity-exchange rates (SCXRs). The average SCXR of 2.78 implies that participants are, on average, indifferent between \$10 for themselves and $\$10 \times 2.78$ for the top-rated charity.[25] Although the main analysis will focus on decisions involving noncensored X_t values, subsequent analyses confirm the robustness to also including decisions that involve censored X_t values.[26]

Each bar shows the percentage of the participants in Study 1 with a given SCXR, which equals $X_t/10$ for each charity t, where participants are indifferent between $\$X_t$ for the top-rated charity t and \$10 for themselves. The results include data for the 97 noncensored X_t values.

Figure 1. Distribution of SCXR in Study 1

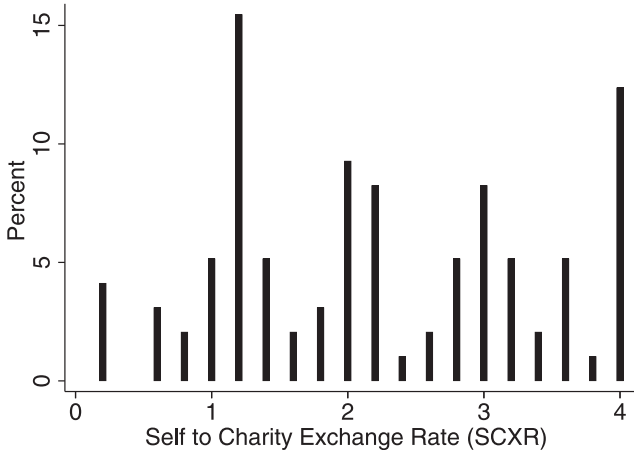

Table 1 presents results from ordinary least squares (OLS) regressions of the valuations for low-rated charities.[27] The coefficient estimate on *charity–self* in column 1 shows that valuations are, on average, a significant 11 percentage points lower in the charity–self treatment than in the charity–charity treatment. In other words, the frequency with which participants are willing to choose money for a lower-rated charity falls by 11 percentage points when their alternative choice involves money for themselves in the charity–self treatment rather than money for the corresponding top-rated charity in the charity–charity treatment. Consistent with excuse-driven responses to performance metrics, participants appear to overweight their dislike of the low-rated charities when self-serving motives are relevant. Indeed, this 11-percentage-point decrease is nearly twice as large as the 5.6-percentage-point decrease seen when going from a second-rated to a third-rated charity (controlled for via rating fixed effects).

Column 2 confirms that this excuse-driven behavior is robust to the inclusion of individual fixed effects and thus not driven by permanent heterogeneity across participants. Column 3 shows that the drop in valuations in the charity–self treatment is more pronounced among individuals who may be more likely to seek excuses not to give: individuals who are more selfish as indicated by higher SCXRs. That excuse-driven responses to charity performance metrics are more likely among more selfish individuals echoes the finding in Exley (2015) that more selfish individuals are more likely to have excuse-driven responses to risk in charitable giving opportunities.[28] Column 4 shows that there are not significant ordering effects between individuals who first complete valuation price lists in the charity–self treatment (captured by the indicator, *order*(*cs*,*cc*)) versus individuals who first complete the valuation price lists in the charity–charity treatment. Such differences may have been expected from a desire to maintain consistency and avoid cognitive dissonance.[29] Column 5 considers variation across charity types. The coefficient on *charity–self* × *KIPP* shows that there is not a significant difference in how participants respond to the lower college matriculation rates of KIPP charter schools versus the lower program expense rates of Make-A-Wish Foundation state chapters (the excluded charity type). Although the positive coefficient on *charity–self* × *animal shelters* implies relatively less evidence for excuse-driven responses to live release rates of animal shelters, the evidence is still significant.[30] Column 6 shows that the results are

Table 1. Ordinary Least Squares Regressions of Valuations for Low-Rated Charities in Study 1

	1	2	3	4	5	6
charity-self	−10.75**	−10.75**	−10.75***	−15.20***	−17.43***	−23.18***
	(4.00)	(4.19)	(3.13)	(4.32)	(5.25)	(4.77)
$(SCXR - \overline{SCXR})$			−8.99***	−9.73***	−9.60***	−9.08***
			(2.94)	(3.21)	(3.30)	(1.68)
charity-self $*(SCXR - \overline{SCXR})$			−12.74***	−11.78***	−11.75***	−13.93***
			(3.40)	(3.42)	(3.44)	(2.18)
order(cs,cc)				−7.53	−7.46	−8.63
				(7.26)	(7.33)	(5.62)
*charity-self * order(cs,cc)*				9.81	9.84	5.11
				(6.06)	(6.13)	(6.00)
KIPP schools					4.54	5.79
					(3.54)	(3.55)
*charity-self * KIPP schools*					0.16	−3.53
					(5.01)	(4.17)
animal shelters					−4.15	−2.77
					(3.02)	(2.52)
*charity-self * animal shelters*					6.77*	7.58**
					(3.95)	(3.54)
Constant	72.22***	35.91***	72.22***	75.63***	75.47***	70.47***
	(3.99)	(2.29)	(3.47)	(4.64)	(4.84)	(4.00)
Rating FEs	Yes	Yes	Yes	Yes	Yes	Yes
Ind FEs	No	Yes	No	No	No	No
Censored X	No	No	No	No	No	Yes
N	388	388	388	388	388	564

Notes. Standard errors are clustered at the individual level and shown in parentheses. The above presents OLS regression results of valuations of a lower-rated charity of type t receiving $\$X_t$. Valuations are scaled as percentages of X_t in the charity–charity treatment and as percentages of \$10 in the *charity–self treatment*. The variable *charity–self* is an indicator for valuations elicited in the charity–self treatment; $(SCXR - \overline{SCXR})$ is an individual's self-to-charity-exchange rate minus the average self-to-charity-exchange rate; *order(cs,cc)* is an indicator for individuals who first complete valuation price lists in the charity-self treatment; and *KIPP schools* and *animal shelters* are indicators for KIPP charter schools and animal shelters, respectively, where the excluded charity type is Make-A-Wish foundation state chapters. "Rating FEs" and "Ind FEs" indicate whether charity rating fixed effects and individual fixed effects are included, respectively. "Censored X" indicates whether cases involving censored X_t values are included. When cases with censored X_t values are not included, the data include valuations from 31 to 35 participants for each charity type. When cases with censored X_t values are included, the data include valuations from 47 participants for each charity type.

*$p < 0.10$; **$p < 0.05$; ***$p < 0.01$.

robust to including valuations involving censored X_t values.[31]

Online Appendix A provides additional evidence for excuse-driven responses to charity performance metrics from two online experiments, Study 1a and Study 1b. Study 1a replicates the main finding of excuse-driven responses to performance metrics on a larger sample of 200 participants using a simplified design. In showing that these results are sensitive to the order in which participants make decisions, Study 1a also provides evidence for the role that a desire to avoid cognitive dissonance and to maintain consistency may play in driving responses. Study 1b, in addition to replicating the main finding of excuse-driven responses on a sample of 201 participants, shows that excuse-driven responses to performance metrics are not mitigated when individuals have an opportunity to exert effort in order to reallocate their giving to a more effective charity.

3. Study 2: An Online Experiment

Study 2 involves data from 400 Amazon Mechanical Turk workers in January 2018 who made five decisions about how much money to keep for themselves or to instead donate to the Make-A-Wish Foundation.[32] In addition to receiving a \$1 completion fee, participants knew that one of their decisions would be randomly selected to count for payment.[33] Relative to Study 1, Study 2 allows for a test of excuse-driven responses to charity performance metrics on a larger sample and via an identification strategy that does not require a normalization procedure. The design and results for Study 2 are detailed below (and see Online Appendix B.4 for instructions and screenshots).

Participants are randomly assigned to the aggregated-information treatment ($n = 201$) or to the disaggregated-information treatment ($n = 199$). In both treatments, participants make five decisions about how much money to keep for themselves, out of an additional payment of 50 cents, versus how much money to instead donate to the Make-A-Wish Foundation. All that varies across their five decisions is the impact of their donation—more specifically, whether the amount they choose to donate is multiplied by 5, 4.5, 4, 3.5, or 3 before being donated to the Make-A-Wish Foundation. All that varies across the two treatments is how the information on the impact of their donation is presented. In the aggregated-information treatment, participants are simply informed that any amount they choose to donate will be multiplied by 5, 4.5, 4, 3.5, or 3. In the disaggregated-information treatments, participants are instead informed that any amount they choose to donate will be multiplied by 5 and then discounted by a "processing fee" of 0%, 10%, 20%, 30%, or 40%—thus implying the same overall multipliers as in the aggregated-information treatment. Unlike in the aggregated-information treatment, however, participants in the disaggregated-information treatment may narrowly bracket the "bad" information on the impact of their donation: the processing fee. The ability to narrowly bracket this bad information may, in turn, facilitate their ability to overweight the extent to which they dislike higher processing fees as an excuse not to give.

Table 2 presents results from OLS regressions of donation behavior on the processing fee. The *fee* variable is set equal to 0 when the multiplier is 5 (or the processing fee is 0%), 0.1 when the multiplier is 4.5 (or the processing fee is 10%), 0.2 when the multiplier is 4 (or the processing fee is 20%), 0.3 when the multiplier is 3.5 (or the processing fee is 30%), and 0.4 when the multiplier is 3 (or the processing fee is 40%).

Column 1 presents results on the donation amount. The coefficient estimate on $fee > 0$ shows that if a donation opportunity involves a processing fee (or, equivalently, a multiplier that is less than 5), the average amount participants give does not significantly change in the aggregated-information treatment. The coefficient estimate on *disagg* also shows that when there is no processing fee, the average amount participants give in the aggregated-information treatment is not significantly different from the average amount participants give in the disaggregated-information treatment. The coefficient estimate on *disagg* × *fee*, however, shows that a processing fee causes the average giving to decrease by a significantly greater amount in the disaggregated-information treatment than it does in the aggregated-information treatment. Although a processing fee causes the average giving to insignificantly decrease by 0.35 cents in

Table 2. Ordinary Least Squares Regressions of Donation Behavior in Study 2

Dependent variable	*donation*		*donation* > 0		*donation*		*donation* > 0	
	1	2	3	4	5	6	7	8
fee > 0	−0.35	−0.34	−0.02	−0.02				
	(0.31)	(0.31)	(0.01)	(0.01)				
disagg.	2.29	2.31	−0.05	−0.05	1.65	1.63	−0.05	−0.06
	(1.67)	(1.67)	(0.05)	(0.05)	(1.61)	(1.61)	(0.05)	(0.05)
disagg. * (*fee* > 0)	−4.52***	−4.55***	−0.12***	−0.12***				
	(0.85)	(0.85)	(0.03)	(0.03)				
fee					−0.36	−0.46	−0.12***	−0.12***
					(0.94)	(0.94)	(0.04)	(0.04)
disagg. * *fee*					−14.89***	−14.79***	−0.45***	−0.44***
					(2.61)	(2.58)	(0.09)	(0.09)
Constant	13.53***	13.15***	0.73***	0.72***	13.33***	12.97***	0.73***	0.73***
	(1.13)	(1.14)	(0.03)	(0.03)	(1.11)	(1.12)	(0.03)	(0.03)
Order FEs	No	Yes	No	Yes	No	Yes	No	Yes
N	2,000	2,000	2,000	2,000	2,000	2,000	2,000	2,000

Notes. Standard errors are clustered at the individual level and shown in parentheses. Columns 1, 2, 5, and 6 present OLS regression results of how much participants donate. Columns 3, 4, 7, and 8 present OLS regression results of an indicator on whether participants donate. The variable $fee > 0$ is an indicator for a processing fee greater than 0 or, equivalently, a multiplier less than 5; *dissagg.* is an indicator for participants in the disaggregated-information treatment; *fee* equals 0, 0.1, 0.2, 0.3, or 0.5 when the processing fee is 0%, 10%, 20%, 30%, or 40%, respectively, in the disaggregated-information treatment or when the multiplier equals 5, 4.5, 4, 3.5, or 3, respectively, in the aggregated-information treatment. "Order FEs" indicate whether fixed effects for the order of each decision are included. The data include five observations from each of the 400 participants.

***$p < 0.01$.

the aggregated-information treatment, a processing fee causes the average giving to significantly decrease by an additional 4.52 cents in the disaggregated-information treatment. Given an average donation amount of 12.59 cents (out of the 50 cents), these decreases in average giving correspond to a decrease of 3% in the aggregated-information treatment versus a 39% decrease in the disaggregated-information treatment. That is, participants appear to overweight their dislike of giving opportunities with lower impact factors when they can easily exploit bad information (i.e., when the processing fee is highlighted in the disaggregated-information treatment) as an excuse not to give.

Column 2 shows that these results are robust to including fixed effects for the order in which decisions are made. Columns 3 and 4 show similar results when considering the likelihood of a participant making a donation. Although a processing fee causes the likelihood of giving to insignificantly decrease by 2 percentage points in the aggregated-information treatment, a processing fee causes the average giving to significantly decrease by 12 additional percentage points in the disaggregated-information treatment. Given a baseline giving rate of 64%, these decreases in the likelihood of giving correspond to a decrease of 3% in the aggregated-information treatment versus a decrease of 22% in the disaggregated-information treatment. Columns 4–8 show similar results when instead considering a continuous measure of the underlying processing fee.

In reflecting on the results from Studies 1 and 2, it is useful to note there are clear trade-offs across these two studies—mostly notably, because Study 1 employs a normalization procedure whereas Study 2 does not. Study 1 examines evidence for excuse-driven responses to charity performance metrics by comparing giving decisions that occur when self-serving motives and thus excuses are relevant (in the charity–self treatment) to ones that occur when self-serving motives and thus excuses are *not* relevant (in the charity–charity treatment). The use of the normalization procedure alleviates concerns that this comparison is merely reflective of giving decisions varying under different stakes. That is, absent a normalization procedure, Figure 1 makes clear that differences across contexts where money for the participants is and is not at stake could simply follow from participants valuing the stakes in each context quite differently.

Study 2 considers a different comparison via a framing manipulation. Study 2 compares giving decisions when excuses are more easily developed (in the disaggregated-information treatment) to when excuses are less easily developed (in the aggregated-information treatment). By making a comparison about whether excuses are more or less easily developed, rather than about whether or not excuses are relevant (as in Study 1), Study 2 always considers contexts where money for the participants is at stake. There is no normalization procedure in Study 2 that allows for a simplified decision environment to follow.[34] However, the lack of a normalization procedure in Study 2 prevents the examination of how giving decisions respond to the framing manipulation in a context where self-serving motives are not relevant but the amount of money at stake is known to be similarly valued as the amount of money at stake in a context where self-serving motives are relevant. Put differently, the results in Study 2 do not allow us to identify to what extent the response to the framing manipulation is an excuse-driven response versus a non-excuse-driven response. The results in Study 2 are indicative of excuse-driven responses to the extent that the framing manipulation does indeed facilitate the development of excuses. With that caution in mind, Online Appendix A.3 provides suggestive evidence of the framing manipulation having a smaller impact when self-serving motives and thus excuses are not relevant.

4. Conclusion

This paper documents how individuals may use charity performance metrics as an excuse not to give. The relation between policy and this novel channel through which individuals exploit factors in a self-serving manner is clear. When considering the benefits of providing performance metrics as a tool to encourage more effective giving, it is important to consider the potential downside of facilitating the development of excuses not to give. How to construct solicitations that balance this tension, and that mitigate the potential downside, is worthy of future work.

A few comments may prove useful to this future work. First, excuses are not relevant when how much to give has been decided and the only question is how to distribute said giving. The provision of performance metrics in these cases may indeed encourage more effective giving absent any concerns related to excuse-driven responses. Second, excuses that involve self-serving views of information may be more likely among more marginal givers such as new potential donors or those who have neither given many times nor very much in the past. Third, providing aggregated information—if positive overall—may limit excuse-driven responses to information by making it more difficult to exploit particular pieces of information.

Finally, it is worth noting that the use of performance metrics is ubiquitous and not unique to the nonprofit sector. Companies compete in terms of performance metrics, and workers are often incentivized and

evaluated in terms of performance metrics. A desire to discount the success of companies in which one does not have a stake or a desire to discriminate against certain job candidates may imply that one's view of related performance metrics is tainted by self-serving motives. A perception that the use of performance metrics is objective and fair may exacerbate the impact of related biases. Thus, managers of both nonprofit and for-profit organizations should carefully consider if and to what extent views of performance metrics may be distorted by self-serving motives.

Endnotes

[1] See http://blog.ted.com/2013/03/11/how-to-pick-the-charity-thats-right-for-you. Accessed February 22, 2019.

[2] See https://www.charitynavigator.org/index.cfm?bay=content.view&cpid=4756#.Ut6o3E8o5dg. Accessed February 22, 2019.

[3] See http://www.givewell.org/giving101. Accessed February 22, 2019.

[4] See https://www.thelifeyoucansave.org/about-us/faq. Accessed February 22, 2019.

[5] See https://www.thelifeyoucansave.org/charity-voices/id/1404/dmis-live-saving-interventions-are-high-impact-low-cost-and-based-on-strong-evidence. Accessed February 22, 2019.

[6] See, for instance, Hwang and Powell (2009), van Iwaarden et al. (2009), and Ebrahim and Rangan (2010).

[7] Fong and Oberholzer-Gee (2011) observe that most individuals are unwilling to learn whether the recipient of a potential donation is disabled or a drug user, even though they are less willing to give to drug users. After also finding that most individuals do not choose to become informed about their giving, Butera and Houser (2016) show how delegating giving allocation decisions to others can lead to more effective giving. Niehaus (2014) models one explanation for this behavior: learning performance information may prevent individuals from maximizing their warm glow by holding the most optimistic beliefs about their impact. Additional work on information avoidance and ask avoidance in the treatment of prosocial behavior include Dana et al. (2006, 2007), Broberg et al. (2007), Oberholzer-Gee and Eichenberger (2008), Larson and Capra (2009), Hamman et al. (2010), Matthey and Regner (2011), Nyborg (2011), DellaVigna et al. (2012), Lazear et al. (2012), Knuttson et al. (2013), Bartling et al. (2014), Feiler (2014), Grossman (2014), van der Weele (2014), van der Weele et al. (2014), Kamdar et al. (2015), Trachtman et al. (2015), Lin et al. (2016), Andreoni et al. (2017), and Grossman and van der Weele (2017). See also Golman et al. (2017) for a review.

[8] Relatedly, Kessler et al. (2019) find that rich and powerful alumni give more when they can indicate their charitable giving priorities—even though this indication was not binding.

[9] See Giving USA Foundation (2015) for data on charitable giving as a percentage of GDP from 1974 to 2014.

[10] Among many other possibilities, learning about better performance metrics may also result in individuals thinking they need to give less in order to signal prosocial tendencies (Butera and Horn 2014).

[11] For instance, in their survey paper, Gino et al. (2016, p. 190), note that "when the context provides sufficient flexibility to allow plausible justification that one can both act egoistically while remaining moral, people seize on such opportunities to prioritize self-interest at the expense of morality."

[12] For surveys on the correspondingly vast and broad related literature, see Merritt et al. (2010), Gino et al. (2016), and Bénabou and Tirole (2016). Some examples of related work also include Snyder et al. (1979), Babcock et al. (1995), Hsee (1996), Konow (2000), Coffman (2011), Linardi and McConnell (2011), Shalvi et al. (2011, 2012), Gino and Ariely (2012), Gino et al. (2013), Falk and Szech (2013), Andreoni and Sanchez (2014), Di Tella et al. (2015), Pittarello et al. (2015), Shalvi et al. (2015), Danilov and Saccardo (2016), Schwardman and van der Weele (2016), and Regner (2018).

[13] To my knowledge, the only other paper to provide evidence for motivated reasoning to payoff information that is absent any uncertainty is Exley and Kessler (2018).

[14] Study 1 was advertised to take one hour and all sessions were completed in less than one hour.

[15] See http://www.charitynavigator.org for information on program expense rates. Accessed February 22, 2019.

[16] See http://www.kipp.org for information on college matriculation rates. Accessed February 22, 2019.

[17] See http://www.maddiesfund.org for information on live release rates. Accessed February 22, 2019.

[18] Buffer price lists involve participants receiving \$5 (instead of \$10) but are otherwise identical to the normalization price lists. They are intended to limit "stickiness" in participants' normalization price list decisions across charity types, but unlike the normalization price lists, they do not influence later experimental parameters.

[19] Even though the top-rated charities do not have perfect performance metrics, participants are provided with background information that describes how the top-rated charities exceed a common metric. For Make-A-Wish Foundation state chapters, participants are informed that the highest program expense rate among any state chapters is 90%. For KIPP charter school locations, participants are informed that the average college matriculation rate among students at KIPP charter schools is 80%. For Bay Area animal shelters, participants are informed that the animal rescue community defines "no-kill" shelters as those with live release rates above 90%.

[20] Participants are aware that there will be subsequent price lists when making their decisions in the normalization price lists. The details of any given price list, however, are only revealed as one progresses throughout the study. There is no evidence that participants expected their decisions in the normalization price list to subsequently influence the amounts in the valuation price lists. If participants could have forecasted this design feature, they may have made decisions that imply higher Xt values so that the amounts of money for charities would be higher in the valuation price lists, and then, if anything, the desire to exploit charity performance metrics as an excuse to keep money for themselves would be muted. Moreover, and encouragingly, our results are robust to the exclusion of participants with censored Xt values (columns 1–5 of Table 1) and to the inclusion of participants with censored Xt values (column 6 of Table 1).

[21] This assumes linearity in money, which is made more reasonable by the small stakes involved.

[22] Similar to charity–charity valuations, charity–self valuations equal the percentage such that participants are indifferent between the lower-rated charity of type t receiving \$$Xt$ and themselves receiving some percentage of \$10.

[23] This is comparable to the 42% observed in Exley (2015) and, more broadly, Engel (2011)'s metastudy finding that 36% of dictators do not give anything to their recipients.

[24] To bias against finding excuse-driven evidence, Xt values are estimated as the upper bound of participants' implied indifference ranges so that they weakly prefer \$$Xt$ for the top-rated charity t over \$10 for themselves.

[25] The average SCXRs range from 2.73 to 2.87 across the three price lists when defined according to their order or when defined according

to the charity type involved. These differences are not statistically significant.

[26] Because stronger evidence for excuse-driven behavior emerges for participants with larger Xt values, initially excluding decisions involving censored Xt values seeks to be conservative. It is not possible to include decisions involving the three participants with unknown Xt values as a result of multiple switch points because the experimental code outputted unreasonable Xt values for them during the experiment.

[27] These valuations are set to equal the midpoint of implied ranges from participants' switch points on the valuation price lists, unless the point estimate of 0% or 100% is implied from a participant never choosing or always choosing $\$Xt$ for the P-rated charity t, respectively. In the 7% of valuation price lists with multiple switch points, valuations are estimated by following prior literature that considers only the first switch points. The results are robust to instead excluding any valuations with multiple switch points.

[28] Other literature with potentially related findings include Exley and Kessler (2018), Karlan and Wood (2017), and Exley and Petrie (2018). For a discussion of about potential heterogeneity in self-serving avoidance of information or giving opportunities, see Nyborg (2011) and Lazear et al. (2012).

[29] Literature documenting a desire to avoid cognitive dissonance includes Babcock et al. (1995), Konow (2000), Haisley and Weber (2010), Gneezy et al. (2016), and Golman et al. (2016).

[30] From column 5 in Table 1, the sum of the coefficients on *charity–self* and *charity–self* × *animal shelters* is significantly different from zero ($p < 0.05$). Follow-up responses suggest this may be driven by some participants thinking that high live release rates are indicative of not needing help; thus either low or high live release rates may serve as excuses. Because such a possibility can confound the results, Studies 1a and 1b (detailed in Online Appendix A) examine more unambiguous metrics and strip away unnecessary contextual details about the involved charities.

[31] Because stronger evidence for excuse-driven behavior emerges for participants with larger Xt values (those with higher self-to-charity-exchange rates), and the largest Xt values arise from censored values, initially excluding valuations involving censored Xt values seeks to be conservative. The results in columns 1–6 are robust to instead considering Tobit regressions.

[32] Study 2 was advertised as a 5–10 minute study, and the median completion time was 4 minutes.

[33] In the taxonomy proposed in Charness et al. (2013), this can be considered an extralaboratory study.

[34] I thank the review team for encouraging this new approach in Study 2 that allowed for a simplified decision environment. It is also worth noting that this new approach neither relies on the assumption of linearity nor could be influenced by moral crediting concerns. Note, however, that the assumption of linearity in Study 1 is reasonable given the small stakes involved and that there is no evidence for moral crediting concerns in Study 1 (i.e., recall that participants who are more selfish in their decisions during the normalization procedure are more likely to be excuse driven and thus more selfish again).

References

Andreoni J, Sanchez AL (2014) Do beliefs justify actions or do actions justify beliefs? An experiment on stated beliefs, revealed beliefs, and social-image motivation. Working paper, University of California, San Diego, La Jolla.

Andreoni J, Rao JM, Trachtman H (2017) Avoiding the ask: A field experiment on altruism, empathy, and charitable giving. *J. Political Econom.* 125(3):625–653.

Babcock L, Loewenstein G, Issacharoff S, Camerer C (1995) Biased Judgments of fairness in bargaining. *Amer. Econom. Rev.* 85(5): 1337–1343.

Bartling B, Engl F, Weber RA (2014) Does willful ignorance deflect punishment?—An experimental study. *Eur. Econom. Rev.* 70(August):512–524.

Batista C, Silverman D, Yang D (2015) Directed giving: Evidence from an inter-household transfer experiment. *J. Econom. Behav. Organ.* 118(October):2–21.

Bénabou R, Tirole J (2016) Mindful economics: The production, consumption, and value of beliefs. *J. Econom. Perspect.* 30(3):141–164.

Broberg T, Ellingsen T, Johannesson M (2007) Is generosity involuntary? *Econom. Lett.* 94(1):32–37.

Brown AL, Meer J, Williams JF (2017) Social distance and quality ratings in charity choice. *J. Behav. Experiment. Econom.* 66(C): 9–15.

Butera L, Horn J (2014) Good news, bad news, and social image: The market for charitable giving. George Mason University Interdisciplinary Center for Economic Science (ICES) working paper, George Mason University, Washington, DC.

Butera L, Houser D (2018) Delegating altruism: Toward an understanding of agency in charitable giving. *J. Econom. Behav. Organ.* 155:99–109.

Charness G, Gneezy U, Kuhn MA (2013) Experimental methods: Extra-laboratory experiments-extending the reach of experimental economics. *J. Econom. Behav. Organ.* 91(July): 93–100.

Coffman LC (2011) Intermediation reduces punishment (and reward). *Amer. Econom. J.: Microeconom.* 3(4):1–30.

Coffman L (2017) Intermediaries in fundraising inhibit quality-driven charitable donations. *Econom. Inquiry* 55(1):409–424.

Dana J, Cain DM, Dawes RM (2006) What you don't know won't hurt me: Costly (but quiet) exit in dictator games. *Organ. Behav. Human Decision Processes* 100(2):193–201.

Dana J, Weber RA, Kuang JX (2007) Exploiting moral wiggle room: Experiments demonstrating an illusory preference for fairness. *Econom. Theory* 33(1):67–80.

Danilov A, Saccardo S (2016) Disguised discrimination. Working paper, University of Cologne, Cologne, Germany.

DellaVigna S, List J, Malmendier U (2012) Testing for altruism and social pressure in charitable giving. *Quart. J. Econom.* 127(1):1–56.

Di Tella R, Perez-Truglia R, Babino A, Sigman M (2015) Conveniently upset: Avoiding altruism by distorting beliefs about others' altruism. *Amer. Econom. Rev.* 105(11):3416–3442.

Ebrahim A, Rangan V K (2010) The limits of nonprofit impact: A contingency framework for measuring social performance. Working paper, Harvard Business School, Boston.

Eckel CC, Grossman PJ (2003) Rebate versus matching: Does how we subsidize charitable contributions matter? *J. Public Econom.* 87(3): 681–701.

Eckel CC, Grossman PJ (2008) Subsidizing charitable contributions: A natural field experiment comparing matching and rebate subsidies. *Experiment. Econom.* 11(3):234–252.

Eckel CC, Herberich D, Meer J (2017) A field experiment on directed giving at a public university. *J. Behav. Experiment. Econom.* 66(C): 66–71.

Engel C (2011) Dictator games: A meta study. *Experiment. Econom.* 14(4):583–610.

Exley CL (2015) Excusing selfishness in charitable giving: The role of risk. *Rev. Econom. Stud.* 83(2):587–628.

Exley CL, Kessler JB (2018) Motivated errors. HBS Working Paper 18-07, Harvard Business School, Boston.

Exley CL, Petrie R (2018) The impact of a surprise donation ask. *J. Public Econom.* 158(February):152–167.

Falk A, Szech N (2013) Organizations, diffused pivotality and Immoral outcomes. IZA Discussion Paper 7442, IZA Institute of Labor Economics, Bonn, Germany.

Feiler L (2014) Testing models of information avoidance with binary choice dictator games. *J. Econom. Psych.* 45(December): 253–267.

Fong CM, Oberholzer-Gee F (2011) Truth in giving: Experimental evidence on the welfare effects of informed giving to the poor. *J. Public Econom.* 95(5–6):436–444.

Gino F, Ariely D (2012) The dark side of creativity: Original thinkers can be more dishonest. *J. Personality Social Psych.* 102(3): 445–459.

Gino F, Ayal S, Ariely D (2013) Self-serving altruism? The lure of unethical actions that benefit others. *J. Econom. Behav. Organ.* 93(September):285–292.

Gino F, Norton MI, Weber RA (2016) Motivated bayesians: Feeling moral while acting egoistically. *J. Econom. Perspect.* 30(3): 189–212.

Giving USA Foundation (2015) Giving USA 2015: The annual report on philanthropy. Technical report, Indiana University Lilly Family School of Philanthropy, Indianapolis.

Gneezy U, Keenan EA, Gneezy A (2014) Avoiding overhead aversion in charity. *Science* 346(6209):632–635.

Gneezy U, Saccardo S, Serra-Garcia M, van Veldhuizen R (2016) Bribing the self. Working paper, University of California, San Diego, La Jolla.

Golman R, Hagmann D, Loewenstein G (2017) Information avoidance. *J. Econom. Literature* 55(1):1–40.

Golman R, Loewenstein G, Moene KO, Zarri L (2016) The preference for belief consonance. *J. Econom. Perspect.* 30(3):165–188.

Gordon TP, Knock CL, Neely DG (2009) The role of rating agencies in the market for charitable contributions: An empirical test. *J. Accounting Public Policy* 28(6):469–484.

Grossman Z (2014) Strategic ignorance and the robustness of social preferences. *Management Sci.* 60(11):2659–2665.

Grossman Z, van der Weele JJ (2017) Self-image and willful ignorance in social decisions. *J. Eur. Econom. Assoc.* 15(1):173–217.

Haisley EC, Weber RA (2010) Self-serving interpretations of ambiguity in other-regarding behavior. *Games Econom. Behav.* 68(2): 614–625.

Hamman JR, Loewenstein G, Weber RA (2010) Self-Interest through delegation: An additional rationale for the principal-agent relationship. *Amer. Econom. Rev.* 100(4):1826–1846.

Hope Consulting (2010) Money for good the US market for impact investments and charitable gifts from individual donors and investors. Technical report, Hope Consulting, San Francisco.

Hsee CK (1996) Elastic justification: How unjustifiable factors influence judgments. *Organ. Behav. Human Decision Processes* 66(1): 122–129.

Huck S, Rasul I (2011) Matched fundraising: Evidence from a natural field experiment. *J. Public Econom.* 95(5–6):351–362.

Hwang H, Powell WW (2009) The rationalization of charity: The influences of professionalism in the nonprofit sector. *Admin. Sci. Quart.* 54(2):268–298.

Kamdar A, Levitt SD, List JA, Mullaney B, Syverson C (2015) Once and done: Leveraging behavioral economics to increase charitable contributions. SPI Working Paper 55, Science of Philanthropy Initiative, University of Chicago, Chicago.

Karlan D, List JA (2007) Does price matter in charitable giving? Evidence from a large-scale natural field experiment. *Amer. Econom. Rev.* 97(5):1774–1793.

Karlan D, Wood DH (2017) The effect of effectiveness: Donor response to aid effectiveness in a direct mail fundraising experiment. *J. Behav. Experiment. Econom.* 66(C):1–8.

Karlan D, List JA, Shafir E (2011) Small matches and charitable giving: Evidence from a natural field experiment. *J. Public Econom.* 95(5): 344–350.

Kessler JB, Milkman KL, Yiwei Zhang C (2019) Getting the rich and powerful to give. *Management Sci.*, ePub ahead of print April 26, https://doi.org/10.1287/mnsc.2018.3142.

Knuttson M, Martinsson P, Wollbrant C (2013) Do people avoid opportunities to donate? A natural field experiment on recycling and charitable giving. *J. Econom. Behav. Organ.* 93(September): 71–77.

Konow J (2000) Fair Shares: Accountability and cognitive dissonance in allocation decisions. *Amer. Econom. Rev.* 90(4):1072–1092.

Larson T, Capra MC (2009) Exploiting moral wiggle room: Illusory preference for fairness? A comment. *Judgment Decision Making* 4(6):467–474.

Lazear EP, Malmendier U, Weber RA (2012) Sorting in experiments with application to social preferences. *Amer. Econom. J.: Appl. Econom.* 4(1):136–163.

Li SX, Eckel C, Grossman PJ, Brown TL (2015) Directed giving enhances voluntary giving to government. *Econom. Lett.* 133(August):51–54.

Lin SC, Schaumberg RL, Reich T (2016) Sidestepping the rock and the hard place: The private avoidance of prosocial requests. *J. Experiment. Soc. Psych.* 64(May):35–40.

Linardi S, McConnell MA (2011) No excuses for good behavior: Volunteering and the social environment. *J. Public Econom.* 95(5–6):445–454.

Matthey A, Regner T (2011) Do I really want to know? A cognitive dissonance-based explanation of other-regarding behavior. *Games* 2(1):114–135.

Meer J (2014) Effects of the price of charitable giving: Evidence from an online crowdfunding platform. *J. Econom. Behav. Organ.* 103(July):113–124.

Meer J (2017) Does fundraising create new giving? *J. Public Econom.* 145(January):82–93.

Meier S (2007) Do subsidies increase charitable giving in the long run? Matching donations in a field experiment. *J. Eur. Econom. Assoc.* 5(6):1203–1222.

Merritt AC, Effron DA, Monin B (2010) Moral self-licensing: When begin good frees us to be bad. *Soc. Personality Psych. Compass* 4(5): 344–357.

Metzger L, Günther I (2019) Making an impact? The relevance of information on aid effectiveness for charitable giving. A laboratory experiment. *J. Dev. Economics*. 136:18–33.

Niehaus P (2014) A theory of good intentions. Working paper, University of California, San Diego, La Jolla.

Null C (2011) Warm glow, information, and inefficient charitable giving. *J. Public Econom.* 95(5–6):455–465.

Nyborg K (2011) "I don't want to hear about it": Rational ignorance among duty-oriented consumers. *J. Econom. Behav. Organ.* 79(3): 263–274.

Oberholzer-Gee F, Eichenberger R (2008) Fairness in extended dictator game experiments. *B.E. J. Econom. Anal. Policy* 8(1):1–21.

Pittarello A, Leib M, Gordon-Hecker T, Shalvi S (2015) Justifications shape ethical blind spots. *Psych. Sci.* 26(6):794–804.

Regner T (2018) Reciprocity under moral wiggle room: Is it a preference or a constraint? *Experiment. Econom.* 21(4):779–792.

Schwardman P, van der Weele JJ (2016) Deception and self-deception. Tinbergen Institute discussion paper, University of Amsterdam, Amsterdam.

Shalvi S, Eldar O, Bereby-Meyer Y (2012) Honesty requires time (and lack of justifications). *Psych. Sci.* 23(10):1264–1270.

Shalvi S, Dana J, Handgraaf MJJ, De Dreu CKW (2011) Justified ethicality: Observing desired counterfactuals modifies ethical

perceptions and behavior. *Organ. Behav. Human Decision Processes* 115(2):181–190.

Shalvi S, Gino F, Barkan R, Ayal S (2015) Self-serving justifications doing wrong and feeling moral. *Current Directions Psych. Sci.* 24(2):125–130.

Snyder ML, Kleck RE, Strenta A, Mentzer SJ (1979) Avoidance of the handicapped: An attributional ambiguity analysis. *J. Personality Social Psych.* 37(12):2297–2306.

Trachtman H, Steinkruger A, Wood M, Wooster A, Andreoni J, Murphy JJ, Rao JM (2015) Fair weather avoidance: Unpacking the costs and benefits of "avoiding the ask." *J. Econom. Sci. Assoc.* 1(1):8–14.

van der Weele JJ (2014) Inconvenient truths: Determinants of strategic ignorance in moral dilemmas. Working paper, University of Amsterdam, Amsterdam.

van der Weele JJ, Kulisa J, Kosfeld M, Friebel G (2014) Resisting moral wiggle room: How robust is reciprocal behavior? *Amer. Econom. J.: Microeconom.* 6(3):256–264.

van Iwaarden J, van der Wiele T, Williams R, Moxham C (2009) Charities: How important is performance to donors? *Internat. J. Quality Reliability Management* 26(1):5–22.

Yörük BK (2016) Charity ratings. *J. Econom. Management Strategy* 25(1):195–219.